

Mechanising Denotational Semantics in Agda

Complete Code Listing

Peter D. Mosses^{a,b,1} Jesper Cockx^{a,2} Bernhard Reus^{c,3}

^a *Department of Software Technology
Delft University of Technology
Delft, Netherlands*

^b *Faculty of Science and Engineering
Swansea University
Swansea, United Kingdom*

^c *School of Engineering and Informatics
University of Sussex
Brighton, United Kingdom*

Abstract

Mechanisation of a mathematical definition (also referred to as formalisation) has many benefits. Here, we focus on mechanising denotational semantic definitions of programming languages by embedding them in the Agda language. The Agda type-checker detects and reports any issues with the wellformedness and type correctness of the embedded definitions.

To minimise the effort required, and to facilitate correlation of the original definition with its Agda embedding, mechanisation should not involve significant reformulation or extension. Here, we show how to embed conventional Scott–Strachey denotational definitions in Agda with almost no changes to λ -notation or domain equations.

Agda notation for definitions of types and functions corresponds closely to the conventional meta-notation of denotational semantics. We have developed a collection of Agda modules with postulated types for commonly used domain constructors and their associated operations. Some of our postulates are inconsistent with a classical set-theoretic interpretation of Agda; we conjecture that they would be consistent with an interpretation of Agda in the higher-order intuitionistic logic used by Simpson in his work on synthetic domain theory.

We illustrate our approach with mechanisations of three denotational definitions: Scott’s D_∞ model of the untyped λ -calculus, Plotkin’s denotational semantics of PCF, and a semantics of a sublanguage of Scheme. In previous work, similar mechanisations in Agda have revealed several unsuspected wellformedness issues in published denotational definitions.

Keywords: Denotational semantics, formalisation, mechanisation, Agda, postulates, synthetic domain theory

¹ Email: p.d.mosses@tudelft.nl

² Email: j.g.h.cockx@tudelft.nl

³ Email: bernhard@sussex.ac.uk

Contents

2	Background	3
2.1	Denotational Semantics	3
2.2	Agda	3
	Notes on some of Agda's syntax and features	4
3	Postulated Domain Notation	4
3.1	Domains	4
3.2	Function Domains	5
3.3	Recursive Domains	5
3.4	Flat Domains	6
3.4.1	Booleans	6
3.4.2	Naturals	6
3.5	Sum Domains	6
3.6	Product Domains	7
3.6.1	Tuples	7
3.6.2	Sequences	7
3.7	Updates	7
4	Illustrative Examples	9
4.1	Untyped Lambda-Calculus	9
4.1.1	Abstract Syntax	9
4.1.2	Domain Equations	9
4.1.3	Semantic Functions	9
4.2	PCF: A Programming Language for Computable Functions	10
4.2.1	Abstract Syntax	10
4.2.2	Domain Equations	10
4.2.3	Semantic functions	11
4.3	<i>Scm</i> : A Sublanguage of <i>Scheme</i>	12
4.3.1	Abstract Syntax	12
4.3.2	Domain Equations	12
4.3.3	Auxiliary Functions	13
4.3.4	Semantic Functions	14
5	Postulated Properties	15
6	Illustrative Tests	15
6.1	LC Tests	16
6.2	PCF Tests	16
	References	17

2 Background

2.1 Denotational Semantics

A conventional Scott–Strachey denotational semantics of a programming language consists of definitions of *abstract syntax*, *semantic domains*, and *semantic functions*. The abstract syntax uses a *context-free grammar* to define sets of abstract syntax trees (ASTs) that model the compositional structure of programs and their phrases. The semantic domains are defined by equations in terms of *domain constructors*, and can be mutually recursive. The semantic functions are defined *inductively*, also with mutual recursion, and map ASTs to their *denotations*, which are elements of domains. The denotation of an AST models its contribution to the semantics of complete programs, and is defined *compositionally* in terms of the denotations of its direct components. The AST arguments of semantic functions are enclosed by $\llbracket \dots \rrbracket$, to avoid potential confusion between abstract syntax terms and semantic notation.

Domains were originally required to be continuous lattices [11,12,14,15]. Various more general notions of domain have been suggested (see [1]) but for us it is only important that endomaps between domains have fixpoints, recursive domain equations have solutions, and each domain comes with a bottom element. The notation for domain constructors and their accompanying operations in a conventional Scott–Strachey semantics, summarised below, is independent of the choice of domains.

Let D, E be domains with elements $\delta : D, \epsilon : E$.

- Every domain D has a *bottom* element $\perp_D$ that represents absence of information (e.g., due to an error or nontermination of a computation), usually written just $\perp$.
- A function $\phi : D \rightarrow E$ is (Scott-) *continuous* if it is monotone and preserves limits of directed subsets of D [1]. The *function domain* $F = D \rightarrow E$ consists of all continuous (total) functions from D to E . The set of all functions from a set A to a domain also forms a domain, ordered pointwise. Functions between domains are defined using abstraction $\lambda\delta.\epsilon$, application $\phi\delta$, the operation *fix* that maps each endofunction $\phi : D \rightarrow D$ to its least fixed point, and the operations associated with other domain constructors.
- For any set A the *flat domain* $A_\perp$ is formed by adding $\perp$ as a fresh element. Functions on sets are implicitly extended to (continuous) functions on flat domains, returning $\perp$ when any argument is $\perp$. Elements τ of the domain of *truth-values* $\mathbf{T} = \{\text{true}, \text{false}\}_\perp$ are used in *conditionals* written $\tau \rightarrow \delta_1, \delta_2$, where $\text{true} \rightarrow \delta_1, \delta_2$ is δ_1 , $\text{false} \rightarrow \delta_1, \delta_2$ is δ_2 , and $\perp_{\mathbf{T}} \rightarrow \delta_1, \delta_2$ is $\perp_D$.
- The *separated sum domain* $X = \dots + Y + \dots$ consists of injected elements written ‘ v in X ’ (where $v : Y$ for some summand Y) together with $\perp_X$. The $\mathbf{T}$ -valued operation $\chi \mathbf{E} Y$ (written $\chi \in Y$ in $\llbracket (\text{Scheme}) \rrbracket$) tests whether $\chi : X$ is the injection of some $v : Y$; if so, $\chi \mid Y$ projects χ to v , otherwise to $\perp_Y$.
- The *product domain* $P = D \times E$ consists of pairs $\langle \delta, \epsilon \rangle$ with $\perp_P = \langle \perp_D, \perp_E \rangle$. When $\pi : P$, the operations $\pi \downarrow 1$ and $\pi \downarrow 2$ select its components; similarly for domains of *n-tuples* D^n and finite *sequences* D^* . Further operations on D^* include the empty sequence $\langle \rangle$, concatenation $\delta_1^* \S \delta_2^*$, length $\# \delta^*$, *n*th component $\delta^* \downarrow n$, and *n*th tail $\delta^* \uparrow n$ ($n \geq 1$).

Robert Tennent’s highly accessible tutorial on denotational semantics [15] includes illustrative examples using the above notation (partly in continuation-passing style) and further details of the conventional Scott–Strachey style.

2.2 Agda

Agda [16,17] is both a strongly typed functional *programming language* with support for first-class *dependent types* and a *proof assistant* based on the Curry–Howard correspondence between propositions and types. A number of features set Agda apart from a typical functional programming language:

- Types in Agda are first-class values of a *universe* Set_i where $\text{Set} = \text{Set}_0$, and each universe Set_i belongs to the next universe Set_{i+1} .
- Agda has *dependent function types* $(x : A) \rightarrow B$ or $\forall x \rightarrow B$ where the type B can depend on the value x .
- All functions in Agda are *total*: evaluating a function call is guaranteed to return a value of the correct type in finite time. To ensure totality, Agda’s type checker includes a termination check for recursive definitions, a positivity check for inductive datatypes, and a consistency check for universe levels.

Thanks to its dependent types and totality, Agda’s type system can be used as a higher-order logic for writing mathematical statements and proofs, where types correspond to propositions and type checking corresponds to checking the validity of the proof.

Notes on some of Agda’s syntax and features

Comments. End-of-line comments are initiated by `--`, while multi-line comments are between `{-` and `-}`.

Naming. Declared symbols, variables, and type constructors all share a single namespace. Names can be any sequence of non-whitespace ASCII and Unicode characters except `.;{}()@`, excluding reserved symbols and keywords such as `→` or `where`. Underscores in names play a special role for infix notation.

Infix notation. Agda supports *infix notation* for defining operators, with underscores in the name of a symbol indicating argument positions. For example, if we declare `_+_ : Nat → Nat → Nat`, we can use it as `1 + 1` (the spaces are required, since `1+1` is a valid Agda name!).

Anonymous functions. Lambda abstractions use the syntax `λ x → u` instead of `λ x. u`.

Implicit arguments. Arguments marked by single curly braces `{...}` in the type of a symbol are considered to be *implicit*. These arguments may be omitted, and are then inferred by Agda’s type checker.

Type classes. Agda has no direct support for type classes, but they can be simulated using Agda’s *instance arguments* to resolve type class instances. Instance arguments are marked by double curly braces `{{...}}` and are resolved automatically by using definitions marked as `instance`.

Rewrite rules. Agda has support for *rewrite rules* [2] that are applied automatically during type checking. Rewrite rules are declared by marking an equality proof `eq : a ≡ b` with a `{-# REWRITE eq #-}` pragma. Agda can optionally check confluence of rewrite rules, but currently does not check their termination. Since only *proven* (or postulated) equalities can be added as rewrite rules, non-confluence and non-termination cannot affect the soundness of Agda’s type checker, only its completeness [3].

3 Postulated Domain Notation

This section postulates Agda notation for the domain constructors and associated functions used in the illustrative examples (Section 4↑). See the accompanying website [4] for hyperlinked, highlighted listings of the complete Agda code with the details elided here (including module imports, fixity declarations, and declarations of the types of meta-variables). In the PDF, the symbol ↑ following a reference to a numbered section is a link to the corresponding page on the website.

3.1 Domains

Domains are embedded in Agda as elements of the type `Domain`. A domain `D` is not itself a type, but it has a *carrier* type `⟦ D ⟧ : Set`, which always contains an element `⊥{D}` (written `⊥` when Agda can infer `D`).

```
module Domains where
  postulate
    Domain : Set                -- Domain is the type of all domains
    ⟦_⟧ : Domain → Set         -- ⟦ D ⟧ is the carrier type of D
    ⊥ : {D : Domain} → ⟦ D ⟧ -- ⊥{D} is the 'bottom' element of D
```

Some previous papers on embedding denotational semantics in Agda [5,6,7] defined domains to be types: `Domain = Set`. However, postulating `⊥ : D` for all `D : Domain` was then *inconsistent* with the existence of an empty type in Agda. Postulating `Domain : Set` avoids that inconsistency.

The notation for domains postulated here supports *type-checking* embeddings of denotational semantics in Agda such as those in the illustrative examples. It does *not* define or constrain the *mathematical structure* of domains, nor the algebraic and universal properties of the associated functions.

3.2 Function Domains

The conventional notation in denotational definitions for the domain of continuous functions from D to E is $D \rightarrow E$ or $[D \rightarrow E]$. However, Agda reserves the notation $D \rightarrow E$ for the *type* of *all* (total) functions from type D to type E ; instead, we use the notation $D \rightarrow^c E$ for embedding continuous function domains:

```
module Functions where
  postulate  $\_ \rightarrow^c \_ : \text{Domain} \rightarrow \text{Domain} \rightarrow \text{Domain}$ 
```

Both λ -abstraction and application preserve continuity. In conventional denotational semantics, functions between domains are defined using λ -abstraction and application from primitive continuous functions associated with specific domain constructors, so they are *automatically* continuous. This motivates treating the carrier $\ll D \rightarrow^c E \gg$ of the embedding of a function domain as a type of continuous functions. (*Proving* functions defined in λ -notation to be continuous in Agda requires pairing each λ -abstraction with an explicit proof of its continuity, which is quite impractical – especially when embedding denotations defined in continuation-passing style.)

However, to support type-checking the *direct* embedding of λ -notation from conventional denotational definitions in Agda, it appears to be necessary to *rewrite* the carrier types of function domains to ordinary function types:

```
postulate dom-cts :  $\ll D \rightarrow^c E \gg \equiv (\ll D \gg \rightarrow \ll E \gg)$ 
{-# REWRITE dom-cts #-}
```

Similarly, the notation $A \rightarrow^s D$ is the embedding of the domain of all functions from an ordinary type A to a domain D (which are trivially continuous, ordered pointwise):

```
postulate  $\_ \rightarrow^s \_ : \text{Set} \rightarrow \text{Domain} \rightarrow \text{Domain}$ 
postulate set-cts :  $\ll A \rightarrow^s D \gg \equiv (A \rightarrow \ll D \gg)$ 
{-# REWRITE set-cts #-}
```

Embeddings of *endofunctions* φ on a domain D should always have fixed points $\text{fix } \varphi$, with fix itself also being continuous:

```
postulate fix :  $\ll (D \rightarrow^c D) \rightarrow^c D \gg$ 
```

3.3 Recursive Domains

Conventional denotational semantics often involves groups of mutually recursive domain definitions. In Agda, recursive type definitions lead to non-termination of the type-checker. To avoid non-termination, it is sufficient to break the recursion by leaving (one or more) domains as *postulated*. The following operations can then be used to map values from a postulated domain to its structure and *vice versa*.

```
module Recursion where
  postulate
     $\_ \cong \_ : \text{Domain} \rightarrow \text{Domain} \rightarrow \text{Set}$ 
    -- an instance of  $D \cong E$  declares that the structure of  $D$  is the same as  $E$ 
    unfold :  $\{\{D \cong E\}\} \rightarrow \ll D \rightarrow^c E \gg$ 
    fold :  $\{\{D \cong E\}\} \rightarrow \ll E \rightarrow^c D \gg$ 
```

The *instance parameter* $\{\{D \cong E\}\}$ of the above operations restricts them to domains D and E for which instance $_ : D \cong E$ has been declared.

3.4 Flat Domains

Lifting an ordinary set A by adding a $\perp$ element gives a flat domain, usually written $A_\perp$. Our Agda embedding postulates a corresponding domain constructor $A + \perp$, together with a function $\uparrow$ for injecting elements of A into $A + \perp$, and an operator $f^\#$ for extending functions on A to arguments in $A + \perp$.

```

module Flat where
postulate
  _+_ : Set → Domain                -- A + ⊥ constructs a flat domain
  ↑    : ⟨⟨ A →s (A + ⊥) ⟩⟩         -- (↑ a) injects a into A + ⊥
  _#   : ⟨⟨ (A →s D) →c (A + ⊥) →c D ⟩⟩ -- f # extends f to map ⊥ to ⊥

```

3.4.1 Booleans

The McCarthy conditional operation $\beta \longrightarrow \delta_1, \delta_2$ extends the usual ternary conditional choice to domains. It returns $\perp$ whenever its first argument is $\perp$.

```

module Booleans where
  Bool⊥ = Bool + ⊥
  _→_ , _ : ⟨⟨ Bool⊥ →c D →c D →c D ⟩⟩ -- β → δ1 , δ2 is conditional choice
  _→_ , _ = (λ b δ1 δ2 → if b then δ1 else δ2) #

```

This module also defines $\text{Eq } A$ for use as an instance parameter, restricting operation definitions to types A such that $_ == _ : A \rightarrow A \rightarrow \text{Bool}$, and postulates a $\text{Bool}\perp$ -valued operation $\delta_1 == \perp \delta_2$ on $A + \perp$.

3.4.2 Naturals

Agda allows decimal notation for natural numbers, as well as unary notation using `zero` and `suc`.

```

module Naturals where
  Nat⊥ = Nat + ⊥

```

3.5 Sum Domains

The separated sum $D + E$ of two domains corresponds to lifting the disjoint union of their carrier sets. The following operations can be used directly for binary sums, and iterated for domains with more than two summands.

```

module Sums where
postulate
  _+_ : Domain → Domain → Domain -- D + E is separated sum
  inj1 : ⟨⟨ D →c (D + E) ⟩⟩      -- inj1 δ is injection from D
  inj2 : ⟨⟨ E →c (D + E) ⟩⟩      -- inj2 ε is injection from E
  [_,_] : ⟨⟨ (D →c F) →c (E →c F) →c ((D + E) →c F) ⟩⟩
  -- [ φ , ψ ] applies φ to arguments in D, and ψ to arguments in E

```

Conventional denotational definitions of programming languages (e.g., in [10]) use domain names instead of numerical indices in operations associated with separated sums. The inherently *dependent* types of the Agda embedding of these operations are as follows.

```

postulate
  _≥_ , _↦_ : Domain → Nat → Domain → Set
  _in⊥_ : ⟨⟨ D ⟩⟩ → (E : Domain) → { { E ≥ n ↦ D } } → ⟨⟨ E ⟩⟩ -- δ in⊥ E injection

```

```

_ | ⊥ _      : ⟨⟨ E ⟩⟩ → (D : Domain) → {E ≳ n ⇨ D} → ⟨⟨ D ⟩⟩      -- ε | ⊥ D  projection
_ ∈ ⊥ _      : ⟨⟨ E ⟩⟩ → (D : Domain) → {E ≳ n ⇨ D} → ⟨⟨ Bool ⊥ ⟩⟩ -- ε ∈ ⊥ D  inspection
    
```

The operations are defined only for D and E where an instance of type $E \gtrsim n \mapsto D$ is declared for some n . Instead of defining the summands D of a separated sum domain E by an equation $E = \dots + D + \dots$, the domain E is merely *postulated*, and each summand is declared separately by `instance _ : E ≳ n ⇨ D` (where n should be a different natural number for each summand).

3.6 Product Domains

The carrier of the binary product $D \times E$ of two domains consists of all pairs (d, e) of elements of D and E with the pair $(\perp\{D\}, \perp\{E\})$ as the bottom element $\perp\{D \times E\}$. Neither the product nor pairing is associative. The following operations can be used directly for binary products, and iterated for products of more than two domains.

```

module Products where
postulate
  _ × _ : Domain → Domain → Domain      -- D × E is cartesian product
  _ , _ : ⟨⟨ D →c E →c (D × E) ⟩⟩      -- (δ , ε) is a pair of elements
  _ ↓1 : ⟨⟨ (D × E) →c D ⟩⟩            -- (δ , ε) ↓1 is δ
  _ ↓2 : ⟨⟨ (D × E) →c E ⟩⟩            -- (δ , ε) ↓2 is ε
    
```

3.6.1 Tuples

The domain $D \hat{\ } n$ of n -tuples of elements of a domain D is conventionally written D^n , but Agda does not support the use of variables as superscripts.

```

module Tuples where
  _ ^ _ : Domain → Nat → Domain      -- D ^ n is the domain of n-tuples (n ≥ 0)
    
```

3.6.2 Sequences

The domain D^* of finite sequences of elements of a domain D is conventionally written D^* .

The following notation for the various operations on sequences was introduced in the early 1970s, and is used in the *Scheme* semantics [10]. (The single angle-brackets $\langle \dots \rangle$ used to form sequences are unrelated to the double angle-brackets $\langle\langle D \rangle\rangle$ used for the carrier of domain D .)

```

module Sequences where
postulate
  _ * : Domain → Domain                -- D * is the finite sequence domain
  ⟨⟩   : ⟨⟨ D * ⟩⟩                      -- ⟨⟩ is the empty sequence
  ⟨ _ ⟩ : ⟨⟨ (D ^ suc n) →c D * ⟩⟩      -- ⟨ δ1 , ... ⟩ is a non-empty sequence
  #    : ⟨⟨ D * →c Nat ⊥ ⟩⟩             -- # δ* is the length of sequence δ*
  _ § _ : ⟨⟨ D * →c D * →c D * ⟩⟩      -- δ*1 § δ*2 is sequence concatenation
  _ ↓ _ : ⟨⟨ D * →c Nat →s D ⟩⟩        -- δ* ↓ n is the nth element
  _ † _ : ⟨⟨ D * →c Nat →s D * ⟩⟩      -- δ* † n is the nth tail
    
```

3.7 Updates

When a type A has an equality operation $_ == _ : A \rightarrow A \rightarrow \text{Bool}$, environments $\rho : \langle\langle A \rightarrow^s D \rangle\rangle$ can be ‘updated’ (i.e., extended or overridden) using the conventional notation $\rho [\delta / a]$, defined as follows.

```

module Updates where
  _[_/_] : {{Eq A}} → ⟨⟨ (A →s D) →c D →c A →s (A →s D) ⟩⟩
  ρ [ δ / a ] = λ a' → if a == a' then δ else ρ a'

```

Similarly for stores $\sigma : \langle\langle (A + \perp) \rightarrow^c D \rangle\rangle$:

```

  _[_/_]⊥ : {{Eq A}} → ⟨⟨ ((A + ⊥) →c D) →c D →c (A + ⊥) →c ((A + ⊥) →c D) ⟩⟩
  σ [ δ / α ]⊥ = λ α' → (α ==⊥ α') → δ , σ α'

```

Defining an operation $\mathbf{m} [x \leftarrow y]$ for extension or overriding of *dependent* maps $\mathbf{m}$ is less straightforward, as it involves an equality test that may return an *equivalence proof*.

4 Illustrative Examples

This section illustrates mechanisation of denotational semantics in Agda with three examples, all using the postulated notation (Section 3 ↑) for domains and their associated operations: the Untyped Lambda-Calculus (Section 4.1 ↑), PCF: A Programming Language for Computable Functions (Section 4.2 ↑), and *Scm*: A Sublanguage of *Scheme*.

4.1 Untyped Lambda-Calculus

This section presents our Agda embedding of a denotational semantics of the untyped λ -calculus.

4.1.1 Abstract Syntax

Abstract syntax is *not* regarded as a domain. In conventional Scott–Strachey style denotational semantics, abstract syntax is generally first-order: terms are finite, totally-defined elements.

A variable is written $x\ n$. The argument n merely distinguishes between variables – it is *not* a De Bruin index. The term constructor `var` below includes variables in terms.

In Agda, mixfix notation requires argument positions `_` to be separated by characters other than spaces. The term constructors for function abstraction and application use the Unicode character `⌊` as a separator.

```
module Examples.LC.Abstract-Syntax where
data Var : Set where
  x : Nat → Var
data Exp : Set where
  var _      : Var → Exp      -- variable reference
  (λ _ ⌊ _ ⌋) : Var → Exp → Exp -- function abstraction
  (⌊ _ ⌋ _ ) : Exp → Exp → Exp -- function application
```

4.1.2 Domain Equations

Simply defining $D_\infty = (D_\infty \rightarrow^c D_\infty)$ would lead to non-termination of the Agda type-checker. Instead, we postulate the domain D_∞ , together with a bijection $D_\infty \cong (D_\infty \rightarrow^c D_\infty)$. This declares `unfold` : $\ll D_\infty \rightarrow^c (D_\infty \rightarrow^c D_\infty) \gg$ and `fold` : $\ll (D_\infty \rightarrow^c D_\infty) \rightarrow^c D_\infty \gg$.

```
module Examples.LC.Domain-Equations where
postulate
  D∞ : Domain                -- corresponds to Scott's domain
  instance eqD∞ : D∞ ≅ (D∞ →c D∞) -- bijection
  Env = Var →s D∞ -- environments
```

Use of the conventional notation $\rho \ [\ \delta \ / \ v \]$ for updating an environment ρ to map v to d requires an equality test for variables, elided here.

4.1.3 Semantic Functions

The semantic equations below correspond closely to those found in textbooks on denotational semantics (e.g., [9]). In larger conventional definitions, `fold` and `unfold` are usually left implicit, but Agda does not support this.

```
module Examples.LC.Semantic-Functions where
[ ] : Exp →  $\ll$  Env →c D∞  $\gg$ 
[ var v ] ρ = ρ v
[ (λ v ⌊ e ⌋) ] ρ = fold ( λ δ → [ e ] (ρ [ δ / v ]) )
[ (⌊ e1 ⌋ e2 ⌋) ] ρ = unfold ( [ e1 ] ρ ) ( [ e2 ] ρ )
```

4.2 PCF: A Programming Language for Computable Functions

PCF and its denotational semantics were originally defined by Dana Scott in 1969 [13] with combinators (S, K) instead of λ -abstraction. Gordon Plotkin subsequently defined a denotational semantics for PCF including λ -abstraction [8]. The Agda modules presented below are an embedding of the latter definition.

PCF is an intrinsically typed language: every well-formed term has a unique type. The following grammar summarises the context-free abstract syntax of types σ, τ and terms M, N with variables α_i^σ ($i \geq 0$) and constants c . In Section 4.2.1 $\uparrow$ we reflect Plotkin’s presentation of PCF more accurately by exploiting Agda’s support for dependent types.

$$\sigma, \tau ::= \iota \mid o \mid (\sigma \rightarrow \tau) \quad (1)$$

$$c ::= tt \mid ff \mid \top \mid \mathbf{Y} \mid k_n \mid (+1) \mid (-1) \mid \mathbf{Z} \quad (2)$$

$$M, N ::= \alpha_i^\sigma \mid c \mid (M N) \mid (\lambda \alpha_i^\sigma M) \quad (3)$$

4.2.1 Abstract Syntax

The following abstract syntax of well-formed PCF terms in Agda uses indexed datatype definitions. PCF function types $\sigma \rightarrow \tau$ are written $\sigma \Rightarrow \tau$, and variables α_i^σ are written $\alpha \text{ i } \sigma$ (where the argument i merely distinguishes between variables – it is *not* a De Bruin index).

```
module Examples.PCF.Abstract-Syntax where
```

```
data Types : Set where
   $\iota$       : Types           -- type terms
  o        : Types           -- individuals
  _  $\Rightarrow$  _ : Types  $\rightarrow$  Types  $\rightarrow$  Types -- truth-values
                                           -- functions

data Vars : Types  $\rightarrow$  Set where
   $\alpha$  : Nat  $\rightarrow$  ( $\sigma$  : Types)  $\rightarrow$  Vars  $\sigma$  -- typed variables
                                           --  $\alpha \text{ i } \sigma$  is a variable of type  $\sigma$ 

data  $\mathcal{L}^A$  : Types  $\rightarrow$  Set where
  tt      :  $\mathcal{L}^A$  o           -- typed constants
  ff      :  $\mathcal{L}^A$  o           -- true
   $\top$      :  $\mathcal{L}^A$  (o  $\Rightarrow$   $\sigma \Rightarrow \sigma \Rightarrow \sigma$ ) -- false
  Y       :  $\mathcal{L}^A$  (( $\sigma \Rightarrow \sigma$ )  $\Rightarrow \sigma$ ) -- conditional
  k       : Nat  $\rightarrow$   $\mathcal{L}^A$   $\iota$     -- fixed point
  ( $\mid +1$ ) :  $\mathcal{L}^A$  ( $\iota \Rightarrow \iota$ ) -- numerals
  ( $\mid -1$ ) :  $\mathcal{L}^A$  ( $\iota \Rightarrow \iota$ ) -- successor
  Z       :  $\mathcal{L}^A$  ( $\iota \Rightarrow$  o) -- predecessor
                                           -- zero test

data Terms : Types  $\rightarrow$  Set where
  V_      : Vars  $\sigma \rightarrow$  Terms  $\sigma$  -- typed terms
  L_      :  $\mathcal{L}^A$   $\sigma \rightarrow$  Terms  $\sigma$  -- variable
  ( $\mid \_ \_ \_$ ) : Terms ( $\sigma \Rightarrow \tau$ )  $\rightarrow$  Terms  $\sigma \rightarrow$  Terms  $\tau$  -- constant
  ( $\mid \lambda \_ \_ \_$ ) : Vars  $\sigma \rightarrow$  Terms  $\tau \rightarrow$  Terms ( $\sigma \Rightarrow \tau$ ) -- function application
                                           -- function abstraction
```

4.2.2 Domain Equations

The domains $\mathcal{D} \sigma$ form a “standard collection of domains for arithmetic” in PCF, written $\mathcal{D}_\sigma$ in [8]. As PCF is a simply-typed language, the domains $\mathcal{D} \sigma$ are not reflexive, so their embedding in Agda can use ordinary type definitions, not involving bijections.

module Examples.PCF.Domain-Equations where

```

 $\mathcal{D}$  : Types → Domain      -- standard domains
 $\mathcal{D} \iota$       = Nat $\perp$         -- natural numbers
 $\mathcal{D} \circ$        = Bool $\perp$       -- truth-values
 $\mathcal{D} (\sigma \Rightarrow \tau) = \mathcal{D} \sigma \rightarrow^c \mathcal{D} \tau$  -- functions
    
```

Environments ρ are type-preserving maps from variables to values. They are naturally modeled by a dependent type: $\text{Env } \sigma$ consists of type-preserving maps from variables in $\text{Vars } \sigma$ to their values in the domain $\mathcal{D} \sigma$. The environment $\rho \perp$ maps all variables to $\perp$.

```

Env = ( $\sigma$  : Types) →  $\langle\langle \text{Vars } \sigma \rightarrow^s \mathcal{D} \sigma \rangle\rangle$  -- typed environments
 $\rho \perp$  : Env                                           -- initial environment
 $\rho \perp \_ \_ = \perp$ 
    
```

Extension or overriding typed environments, written $\rho [v / x]'$, requires instances of the equality tests for both variables and types. The definition of the latter is somewhat tedious.

4.2.3 Semantic functions

The notation $\rho \llbracket \alpha \text{ i } \sigma \rrbracket$ gives the value of the variable $\alpha \text{ i } \sigma$ in ρ by applying $\rho \sigma$ to the variable.

module Examples.PCF.Semantic-Functions where

```

 $\_ \llbracket \_ \rrbracket$  : Env → Vars  $\sigma$  →  $\langle\langle \mathcal{D} \sigma \rangle\rangle$       -- typed variable denotations
 $\rho \llbracket \alpha \text{ i } \sigma \rrbracket = \rho \sigma (\alpha \text{ i } \sigma)$ 
    
```

The semantic function $\mathcal{A} \llbracket c \rrbracket$ gives the standard interpretation of the constant c . The corresponding definitions in [8] use case analysis on the domain $\mathcal{D} \iota$, which our Agda embedding does not support (partly because it can express non-continuous functions).

```

 $\mathcal{A} \llbracket \_ \rrbracket$  :  $\mathcal{L}^A \sigma \rightarrow \langle\langle \mathcal{D} \sigma \rangle\rangle$       -- typed constant denotations
 $\mathcal{A} \llbracket \text{tt} \rrbracket$       =  $\uparrow \text{true}$ 
 $\mathcal{A} \llbracket \text{ff} \rrbracket$       =  $\uparrow \text{false}$ 
 $\mathcal{A} \llbracket \supset \rrbracket$       =  $\lambda \beta \delta_1 \delta_2 \rightarrow (\beta \longrightarrow \delta_1, \delta_2)$ 
 $\mathcal{A} \llbracket \text{Y} \rrbracket$        = fix
 $\mathcal{A} \llbracket \text{k n} \rrbracket$      =  $\uparrow n$ 
 $\mathcal{A} \llbracket (+1) \rrbracket$     =  $(\lambda n \rightarrow \uparrow (n + 1))^\#$ 
 $\mathcal{A} \llbracket (-1) \rrbracket$     =  $(\lambda n \rightarrow \uparrow (n ==^N 0) \longrightarrow \perp, \uparrow (n - 1))^\#$ 
 $\mathcal{A} \llbracket Z \rrbracket$        =  $(\lambda n \rightarrow \uparrow (n ==^N 0))^\#$ 
    
```

The semantic function $\mathcal{A}' \llbracket M \rrbracket$ is written $\hat{\mathcal{A}} \llbracket M \rrbracket$ in [8]. It gives the denotation of the term M as a function of the environment ρ .

```

 $\mathcal{A}' \llbracket \_ \rrbracket$  : Terms  $\sigma \rightarrow \langle\langle \text{Env} \rightarrow^s \mathcal{D} \sigma \rangle\rangle$  -- typed term denotations
 $\mathcal{A}' \llbracket V \alpha \text{ i } \sigma \rrbracket \rho$       =  $\rho \llbracket \alpha \text{ i } \sigma \rrbracket$ 
 $\mathcal{A}' \llbracket L c \rrbracket \rho$              =  $\mathcal{A} \llbracket c \rrbracket$ 
 $\mathcal{A}' \llbracket (\text{M } \sqcup \text{N}) \rrbracket \rho$       =  $\mathcal{A}' \llbracket M \rrbracket \rho (\mathcal{A}' \llbracket N \rrbracket \rho)$ 
 $\mathcal{A}' \llbracket (\lambda \alpha \text{ i } \sigma \sqcup M) \rrbracket \rho x$  =  $\mathcal{A}' \llbracket M \rrbracket (\rho [x / \alpha \text{ i } \sigma]')$ 
    
```

Comparison with Plotkin's original definition of PCF [8] confirms the directness of our Agda embedding.


```

P = L × L           -- pairs
U = Ide →s L         -- environments
data Misc : Set where null unallocated undefined unspecified : Misc
M = Misc + ⊥         -- miscellaneous
postulate E : Domain -- expressed values
S = L →c E           -- stores
postulate A : Domain -- answers
C = S →c A           -- command continuations
F = E * →c (E →c C) →c C -- procedure values
    
```

The published denotational semantics of *Scm* [6] defines the domain **E** of expression values by the equation $\mathbf{E} = \mathbf{T} + \mathbf{R} + \mathbf{P} + \mathbf{M} + \mathbf{F}$, and the domain **F** of procedure values by $\mathbf{F} = \mathbf{E}^* \rightarrow (\mathbf{E} \rightarrow \mathbf{C}) \rightarrow \mathbf{C}$. Mutually recursive groups of domain equations have well-defined solutions, but in Agda, defining both the corresponding domains **E** and **F** by equations would cause the type checker to diverge. Postulating one (or both) of these domains avoids divergence; postulating **E** also has the benefit that the embeddings and projections for its summands subsume the bijection between **E** and its intended structure.

The following postulates instantiate injection ($\delta \text{ in } \perp \mathbf{E}$), inspection ($\varepsilon \in \perp \mathbf{D}$), and projection ($\varepsilon \mid \perp \mathbf{D}$) for each summand **D** of **E**.

```

postulate instance
  E+=T : E ≳ 1 ↦ T
  E+=R : E ≳ 2 ↦ R
  E+=P : E ≳ 3 ↦ P
  E+=M : E ≳ 4 ↦ M
  E+=F : E ≳ 5 ↦ F
    
```

4.3.3 Auxiliary Functions

The λ -notation in the Agda definitions of auxiliary functions for *Scm* corresponds closely to that in its published denotational semantics [6].

```

module Examples.Scm.Auxiliary-Functions where

assign : ⟨ L →c E →c C →c C ⟩ -- assign α ∈ stores ε at location α
assign α ε θ σ = θ (σ [ ε / α ] ⊥)

hold : ⟨ L →c (E →c C) →c C ⟩ -- hold α gives the value stored at α
hold α κ σ = κ (σ α) σ
    
```

In the continuation-passing style used to define auxiliary functions for *Scm*, giving explicit continuity proofs would be particularly tedious. For example, the function `hold` is simply a combination of λ -abstraction and application, which is wellknown to ensure continuity.

```

postulate new : ⟨ (L →c C) →c C ⟩ -- new gives an unallocated location

alloc : ⟨ E →c (L →c C) →c C ⟩ -- alloc ε allocates a location for ε
alloc ε κ = new (λ α → assign α ε (κ α))

postulate initial-store : ⟨ S ⟩ -- may have initialised locations
    
```

Conventional denotational definitions usually leave the injection function $\uparrow$ from sets into flat domains implicit, in contrast to the embedding of the definition of `truish`:

```

truish : ⟨ E →c T ⟩ -- truish ε is true for all ε except false
truish ε = (ε ⊥ T) → (((ε ⊥ T) == ⊥ ↑ false) → ↑ false , ↑ true) ,
           ↑ true

```

The remaining auxiliary function definitions shown here involve the operations for (finite) sequences ϵ^* declared in the module `Notation.Products.Sequences`.

```

cons : ⟨ F ⟩ -- cons ⟨ ε1 , ε2 ⟩ allocates and initialises a pair
cons ε* κ = (# ε* == ⊥ ↑ 2) →
            alloc (ε* ↓ 1) (λ α1 →
                alloc (ε* ↓ 2) (λ α2 → κ ((α1 , α2) in ⊥ E))) ,
            ⊥

```

In [6] the auxiliary function *list* is defined by recursion on ϵ^* . Agda accepts recursive definitions only when it can mechanically prove that the recursion terminates, which is not the case for arguments in postulated types. The following definition uses the postulated operation *fix* to avoid recursion.

```

list : ⟨ F ⟩ -- list ε* allocates and initialises a list
list = fix λ (list' : ⟨ F ⟩) → λ ε* κ →
        (# ε* == ⊥ ↑ 0) → κ (↑ null in ⊥ E) ,
        list' (ε* ↑ 1) (λ ε → cons ⟨ (ε* ↓ 1) , ε ⟩ κ)

```

4.3.4 Semantic Functions

The λ -notation in the Agda definitions of semantic functions for *Scm* corresponds closely to that in its conventional denotational semantics [6].

module `Examples.Scm.Semantic-Functions` where

```

K[ ] : ⟨ Con →s E ⟩ -- constant denotations
E[ ] : ⟨ Exp →s U →c (E →c C) →c C ⟩ -- expression denotations
E*[ ] : ⟨ Exp* →s U →c (E* →c C) →c C ⟩ -- sequence denotations

K[ int Z ] = ↑ Z in ⊥ E
K[ #t ] = ↑ true in ⊥ E
K[ #f ] = ↑ false in ⊥ E

E[ con K ] ρ κ = κ (K[ K ])
E[ ide l ] ρ κ = hold (ρ l) κ
E[ (E ⊔ E*) ] ρ κ = E[ E ] ρ (λ ε → E*[ E* ] ρ (λ ε* → (ε ⊥ F) ε* κ))
E[ (lambda l ⊔ E) ] ρ κ = κ ( (λ ε* κ' →
                                list ε* (λ ε → alloc ε (λ α → E[ E ] (ρ [ α / l ]) κ'))
                                ) in ⊥ E )
E[ (if E ⊔ E1 ⊔ E2) ] ρ κ = E[ E ] ρ (λ ε → truish ε → E[ E1 ] ρ κ , E[ E2 ] ρ κ)
E[ (set! l ⊔ E) ] ρ κ = E[ E ] ρ (λ ε → assign (ρ l) ε (κ (↑ unspecified in ⊥ E)))

E*[ ⊔⊔⊔ ] ρ κ = κ ⟨ ⟩
E*[ E ⊔ E* ] ρ κ = E[ E ] ρ (λ ε → E*[ E* ] ρ (λ ε* → κ ((ε) § ε*)))

```

5 Postulated Properties

The `Properties` module postulates basic properties of some of the operations of the postulated domain notation (Section 3↑). These properties are expected to hold in various categories of domains [1] but they do *not* define the *mathematical structure* of domains.

The postulated properties support proofs that terms have identical denotations. For example, some illustrative tests (Section 6↑) declare that the denotation of a function application is equivalent to the denotation of a constant; other tests declare that particular instances of renaming do not affect denotations.

When postulated properties are declared as *rewrite rules*, Agda can use them *automatically* in proofs. Agda also has an option to check that the declared rewrite rules form a confluent system. Rewrite rules are safe to use with `Agda.Builtin.Equality` when that option is enabled. Confluent but non-terminating rewrite rules cannot break consistency, as shown by Cockx, Tabareau, and Winterhalter [(10.1145/3434341)].

The rewrite rules declared below support *automatic* proof of identity for all the illustrative tests: the proof terms are simply `refl` (i.e., reflexivity).

```
module Functions where
postulate
  apply-fix : {φ : ⟨ D →c D ⟩} → fix φ ≡ φ (fix φ) -- apply-fix{φ} unfolds fix φ once
  {-# REWRITE apply-fix #-}
```

The rewrite rule `apply-fix` does not cause the type-checker to diverge, despite the obvious non-termination. Agda's type checker uses *weak head evaluation*: it only unfolds expressions to the point where the top-level constructor becomes visible. In particular, it will not evaluate under a λ -abstraction unless it is being compared to another λ -abstraction and the bodies are not syntactically equal.

```
module Recursion where
postulate
  elim-unfold-fold : { { _ : D ≅ E } } → { e : ⟨ E ⟩ } → unfold (fold e) ≡ e
  {-# REWRITE elim-unfold-fold #-}
```

A rule for `fold (unfold d) ≡ d` could be added, but it is not needed for the current illustrative tests.

```
module Flat where
postulate
  elim-#-↑ : (f #) (↑ a') ≡ f a'
  elim-#-⊥ : (f #) ⊥ ≡ ⊥
  {-# REWRITE elim-#-↑ elim-#-⊥ #-}
```

Removing any of the above rewrite rules breaks the proof in at least one of the illustrative tests. In principle, all `refl` proof terms that rely on rewrite rules could be replaced by proofs that apply the postulated properties to specified subterms. However, we expect that it would be quite tedious to develop such proofs, and reading them is unlikely to provide new insights.

The remaining postulated properties are for domains that are not used in the semantics of the LC and PCF languages; they will be needed when tests for equivalence of denotations of *Scm* expressions are added. Postulated properties for operations on tuples and sequences have not yet been developed.

6 Illustrative Tests

The tests below illustrate *automatic* proof of equivalence of term denotations by the Agda type-checker: the type-correctness of the `refl` definitions of the declared equivalences confirms that they hold. Such tests can check that the denotation of a function call is equivalent to the denotation of a constant in the embedding of a denotational semantics. They can even check that a specific renaming preserves the denotation of a term.

Apart from confirming that a denotational semantics defines denotations which satisfy some expected equivalences at least for some terms, some of the tests also depend on rewrite rules for the postulated

properties (Section 5[↑]). The success of those tests indirectly checks that the rewrite rules preserve denotations. (A more systematic approach would be to develop a suite of unit tests for consequences of postulated properties, independently of denotational definitions.)

6.1 LC Tests

```

check-convergence : -- (λx1.x42)((λx0.x0 x0)(λx0.x0 x0)) = x42
  [ ( (λ x 1 ⊔ var x 42) ⊔
    ( (λ x 0 ⊔ (var x 0 ⊔ var x 0) ) ⊔ (λ x 0 ⊔ (var x 0 ⊔ var x 0) ) ) ) ] ≡ [ var x 42 ]
check-convergence = refl

check-abs : -- (λx1.x1)(λx1.x42) = λx2.x42
  [ ( (λ x 1 ⊔ var x 1) ⊔ (λ x 1 ⊔ var x 42) ) ] ≡ [ (λ x 1 ⊔ var x 42) ]
check-abs = refl

check-free : -- (λx1.(λx42.x1)x2)x42 = x42
  [ ( (λ x 1 ⊔ ( (λ x 42 ⊔ var x 1) ⊔ var x 2 ) ) ⊔ var x 42 ) ] ≡ [ var x 42 ]
check-free = refl
    
```

A reviewer pointed out that the only interpretation of $\ll D_\infty \gg$ in Agda could be a singleton type, so that one should expect many equalities to hold. However, when the proof of an equality is simply by `refl`, Agda's type-checker does not automatically use such reasoning.

6.2 PCF Tests

```

check-fix-lambda : -- fix (λg. λa. 42) 2 ≡ 42
  A'[ ( (L Y ⊔ (λ g ⊔ (λ a ⊔ L k 42) ) ) ⊔ L k 2 ) ] ρ⊥ ≡ ↑ 42
check-fix-lambda = refl

check-countdown : -- fix (λg. λa. ifz a then 42 else g (pred a)) 5 ≡ 42
  A'[ ( (L Y ⊔ (λ g ⊔ (λ a ⊔
    ( (L ⊃ ⊔ (L Z ⊔ V a) ) ⊔ L k 42 ) ⊔
    ( V g ⊔ (L (−1) ⊔ V a) ) ) ) ) ⊔ L k 5 ) ] ρ⊥ ≡ ↑ 42
check-countdown = refl
    
```


References

- [1] Abramsky, S. and A. Jung, *Domain theory*, pages 1–168, Oxford University Press, Inc., USA (1995), ISBN 019853762X.
<https://achimjungbham.github.io/pub/papers/handy1.pdf>
- [2] Cockx, J., *Type theory unchained: Extending Agda with user-defined rewrite rules*, in: M. Bezem and A. Mahboubi, editors, *25th International Conference on Types for Proofs and Programs (TYPES 2019)*, volume 175 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 2:1–2:27, Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany (2020), ISBN 978-3-95977-158-0, ISSN 1868-8969.
<https://doi.org/10.4230/LIPIcs.TYPES.2019.2>
- [3] Cockx, J., N. Tabareau and T. Winterhalter, *The taming of the rew: a type theory with computational assumptions*, Proc. ACM Program. Lang. **5** (2021).
<https://doi.org/10.1145/3434341>
- [4] *Mechanising denotational semantics in Agda: Code* (2026).
<https://pdmosses.github.io/mfps2026-agda/>
- [5] Mosses, P. D., *Checking a denotational semantics of Scheme in Agda*, in: *Proceedings of the 26th ACM SIGPLAN International Workshop on Scheme and Functional Programming, Scheme '25*, pages 3–15, Association for Computing Machinery, New York, NY, USA (2025), ISBN 9798400721625.
<https://doi.org/10.1145/3759537.3762694>
- [6] Mosses, P. D., *A compositional semantics for `eval` in Scheme*, in: *Proceedings of the Workshop Dedicated to Olivier Danvy on the Occasion of His 64th Birthday, OLIVIERFEST '25*, pages 72–81, Association for Computing Machinery, New York, NY, USA (2025), ISBN 9798400721502.
<https://doi.org/10.1145/3759427.3760369>
- [7] Mosses, P. D., *Lightweight Agda formalization of denotational semantics*, in: F. Nordvall Forsberg, editor, *Proceedings of the 31st International Conference on Types for Proofs and Programs, TYPES 2025*, pages 286–290, University of Strathclyde, Glasgow, Scotland (2025).
https://msp.cis.strath.ac.uk/types2025/abstracts/TYPES2025_paper11.pdf
- [8] Plotkin, G. D., *LCF considered as a programming language*, Theoretical Computer Science **5**, pages 223–255 (1977), ISSN 0304-3975.
[https://doi.org/10.1016/0304-3975\(77\)90044-5](https://doi.org/10.1016/0304-3975(77)90044-5)
- [9] Reynolds, J. C., *Theories of Programming Languages*, Cambridge Univ. Press, Cambridge, UK (1998), ISBN 9780521106979.
<https://doi.org/10.1017/CB09780511626364>
- [10] *Scheme standards*.
<https://standards.scheme.org>
- [11] Scott, D., *Outline of a mathematical theory of computation*, in: *Proc. Fourth Annual Princeton Conference on Information Sciences and Systems*, pages 169–176, IEEE, Princeton, NJ, USA (1970). Also: Tech. Monograph PRG-2, Oxford Univ. Computing Lab., Programming Research Group (1970). URL <https://www.cs.ox.ac.uk/files/3222/PRG02.pdf>.
<https://ncatlab.org/nlab/files/Scott-TheoryOfComputation.pdf>
- [12] Scott, D., *Continuous lattices*, in: F. W. Lawvere, editor, *Toposes, Algebraic Geometry and Logic*, pages 97–136, Springer Berlin Heidelberg, Berlin, Heidelberg (1972), ISBN 978-3-540-37609-5.
<https://doi.org/10.1007/BFb0073967>
- [13] Scott, D. S., *A type-theoretical alternative to ISWIM, CUCH, OWHY*, Theoretical Computer Science **121**, pages 411–440 (1993), ISSN 0304-3975.
[https://doi.org/10.1016/0304-3975\(93\)90095-B](https://doi.org/10.1016/0304-3975(93)90095-B)
- [14] Stoy, J. E., *Denotational Semantics: The Scott-Strachey Approach to Programming Language Semantics*, MIT Press, Cambridge, MA, USA (1977), ISBN 978-0-262-19147-0.
- [15] Tennent, R. D., *The denotational semantics of programming languages*, Commun. ACM **19**, pages 437–453 (1976), ISSN 0001-0782.
<https://doi.org/10.1145/360303.360308>
- [16] The Agda Team, *Agda Language Reference* (2025).
<https://agda.readthedocs.io/en/v2.8.0/language/>
- [17] Wikipedia, *Agda* (2026).
[https://en.wikipedia.org/wiki/Agda_\(programming_language\)](https://en.wikipedia.org/wiki/Agda_(programming_language))